

Telecom Customer Churn Analysis

1. Domain Exploration

- understand the business process
- Identify beliefs, loopholes, losses, data flow in the business flow

1. Data Collection & data exploration

- Collecting data from different business verticals, preparing a dataset out of it
- explore data for common challenges and data quality issues.

1. Data Cleaning

- Handling missing values
- Handling unwanted columns - identifiers
- Handling duplicate entries
- Handling outliers

1. Descriptive & Exploratory Analysis

- Statistics
- Data Visualization

1. Preparing the report

Data Exploration

In [2]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

In [3]:

```
#Load data
df = pd.read_excel(r"E:\MLIoT\ML\dataset\telecom\telecom_churn_modelling.xlsx")
df.shape
```

Out[3]:

```
(3333, 20)
```

In [4]:

```
df.head()
```

Out[4]:

	State	Account length	Area code	International plan	Voice mail plan	Number vmail messages	Total day minutes	Total day calls	Total day charge	Total eve minutes	
0	KS	128	415	No	Yes	25	265.1	110	45.07	197.4	
1	OH	107	415	No	Yes	26	161.6	123	27.47	195.5	
2	NJ	137	415	No	No	0	243.4	114	41.38	121.2	
3	OH	84	408	Yes	No	0	299.4	71	50.90	61.9	
4	OK	75	415	Yes	No	0	166.7	113	28.34	148.3	

In [5]:

```
len(df['State'].unique())
```

Out[5]:

51

In [6]:

```
df['Area code'].unique()
```

Out[6]:

```
array([415, 408, 510], dtype=int64)
```

Observations -

- unwanted columns - State may be an identifier and may not be needed, Area code
- Customer having no voice mail plan will have the value of number of vmail messages as 0
- mostly total xxx minutes should be highly correlated to total xxx charge

In [7]:

df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3333 entries, 0 to 3332
Data columns (total 20 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   State                                3333 non-null   object
1   Account length                      3333 non-null   int64
2   Area code                          3333 non-null   int64
3   International plan                  3333 non-null   object
4   Voice mail plan                    3333 non-null   object
5   Number vmail messages              3333 non-null   int64
6   Total day minutes                  3333 non-null   float64
7   Total day calls                    3333 non-null   int64
8   Total day charge                   3333 non-null   float64
9   Total eve minutes                  3333 non-null   float64
10  Total eve calls                    3333 non-null   int64
11  Total eve charge                   3333 non-null   float64
12  Total night minutes                3333 non-null   float64
13  Total night calls                  3333 non-null   int64
14  Total night charge                 3333 non-null   float64
15  Total intl minutes                 3333 non-null   float64
16  Total intl calls                   3333 non-null   int64
17  Total intl charge                  3333 non-null   float64
18  Customer service calls             3333 non-null   int64
19  Churn                              3333 non-null   bool
dtypes: bool(1), float64(8), int64(8), object(3)
memory usage: 498.1+ KB
```

In [8]:

df['State'].unique()

Out[8]:

```
array(['KS', 'OH', 'NJ', 'OK', 'AL', 'MA', 'MO', 'LA', 'WV', 'IN', 'RI',
       'IA', 'MT', 'NY', 'ID', 'VT', 'VA', 'TX', 'FL', 'CO', 'AZ', 'SC',
       'NE', 'WY', 'HI', 'IL', 'NH', 'GA', 'AK', 'MD', 'AR', 'WI', 'OR',
       'MI', 'DE', 'UT', 'CA', 'MN', 'SD', 'NC', 'WA', 'NM', 'NV', 'DC',
       'KY', 'ME', 'MS', 'TN', 'PA', 'CT', 'ND'], dtype=object)
```

In [9]:

df['International plan'].unique()

Out[9]:

```
array(['No', 'Yes'], dtype=object)
```

In [10]:

df['Voice mail plan'].unique()

Out[10]:

```
array(['Yes', 'No'], dtype=object)
```

Data Cleaning

In [11]:

```
# check for duplicated rows  
df.duplicated().sum()
```

Out[11]:

0

In [12]:

```
# check for missng values  
df.isnull().sum()
```

Out[12]:

State	0
Account length	0
Area code	0
International plan	0
Voice mail plan	0
Number vmail messages	0
Total day minutes	0
Total day calls	0
Total day charge	0
Total eve minutes	0
Total eve calls	0
Total eve charge	0
Total night minutes	0
Total night calls	0
Total night charge	0
Total intl minutes	0
Total intl calls	0
Total intl charge	0
Customer service calls	0
Churn	0

dtype: int64

In [13]:

```
# Dropping unwanted columns  
# - analyse the categorical columns
```

In [14]:

```
# check for outliers
df.skew()
```

Out[14]:

```
Account length      0.096606
Area code           1.126823
Number vmail messages 1.264824
Total day minutes   -0.029077
Total day calls      -0.111787
Total day charge     -0.029083
Total eve minutes    -0.023877
Total eve calls       -0.055563
Total eve charge     -0.023858
Total night minutes  0.008921
Total night calls     0.032500
Total night charge    0.008886
Total intl minutes   -0.245136
Total intl calls      1.321478
Total intl charge     -0.245287
Customer service calls 1.091359
Churn                2.018356
dtype: float64
```

Descriptive Analysis

In [15]:

```
df.head()
```

Out[15]:

	State	Account length	Area code	International plan	Voice mail plan	Number vmail messages	Total day minutes	Total day calls	Total day charge	Total eve minutes	Total eve charge
0	KS	128	415	No	Yes	25	265.1	110	45.07	197.4	19.07
1	OH	107	415	No	Yes	26	161.6	123	27.47	195.5	19.07
2	NJ	137	415	No	No	0	243.4	114	41.38	121.2	12.12
3	OH	84	408	Yes	No	0	299.4	71	50.90	61.9	6.19
4	OK	75	415	Yes	No	0	166.7	113	28.34	148.3	14.83

State

In [16]:

```
df['State'].value_counts()
```

Out[16]:

WV	106
MN	84
NY	83
AL	80
WI	78
OH	78
OR	78
WY	77
VA	77
CT	74
VT	73
ID	73
MI	73
UT	72
TX	72
IN	71
KS	70
MD	70
MT	68
NJ	68
NC	68
WA	66
NV	66
CO	66
MA	65
MS	65
RI	65
AZ	64
MO	63
FL	63
ND	62
ME	62
NM	62
NE	61
OK	61
DE	61
SD	60
SC	60
KY	59
IL	58
NH	56
AR	55
GA	54
DC	54
TN	53
HI	53
AK	52
LA	51
PA	45
IA	44
CA	34


Name: State, dtype: int64

In [19]:

```
df['State'].value_counts().plot(kind='barh',figsize=(20,10))
```

Out[19]:

<matplotlib.axes._subplots.AxesSubplot at 0x2a4cecb4288>



Observation -

- the number of customers from each state is differently distributed

In [21]:

```
df.columns
```

Out[21]:

```
Index(['State', 'Account length', 'Area code', 'International plan',
      'Voice mail plan', 'Number vmail messages', 'Total day minutes',
      'Total day calls', 'Total day charge', 'Total eve minutes',
      'Total eve calls', 'Total eve charge', 'Total night minutes',
      'Total night calls', 'Total night charge', 'Total intl minutes',
      'Total intl calls', 'Total intl charge', 'Customer service calls',
      'Churn'],
      dtype='object')
```

In [22]:


```
cats = ['Area code', 'International plan', 'Voice mail plan', 'Churn']
for col in cats:
    print(df[col].value_counts())

plt.figure(figsize=(8,4))
sns.countplot(df[col])
plt.show()
```




```
415    1655
510     840
408     838
```

Name: Area code, dtype: int64




```
No    3010
Yes     323
```

Name: International plan, dtype: int64




```
No    2411
Yes     922
```

Name: Voice mail plan, dtype: int64



```
False    2850
True      483
Name: Churn, dtype: int64
```



Observations -

- Area code - almost half of customers are from area code 415, 1/4 from other two area code each
- International Plan - almost 90% of customers do not have international plan
- Voice mail Message - almost 30% of customers have opted for voice mail message
- Churn - almost 14% of customers left the telecom company

Numeric variables

In [23]:

df.describe()

Out[23]:

	Account length	Area code	Number vmail messages	Total day minutes	Total day calls	Total day charge	To rr
count	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.000000	3333.0
mean	101.064806	437.182418	8.099010	179.775098	100.435644	30.562307	200.9
std	39.822106	42.371290	13.688365	54.467389	20.069084	9.259435	50.1
min	1.000000	408.000000	0.000000	0.000000	0.000000	0.000000	0.0
25%	74.000000	408.000000	0.000000	143.700000	87.000000	24.430000	166.0
50%	101.000000	415.000000	0.000000	179.400000	101.000000	30.500000	201.4
75%	127.000000	510.000000	20.000000	216.400000	114.000000	36.790000	235.1
max	243.000000	510.000000	51.000000	350.800000	165.000000	59.640000	363.1

In [24]:

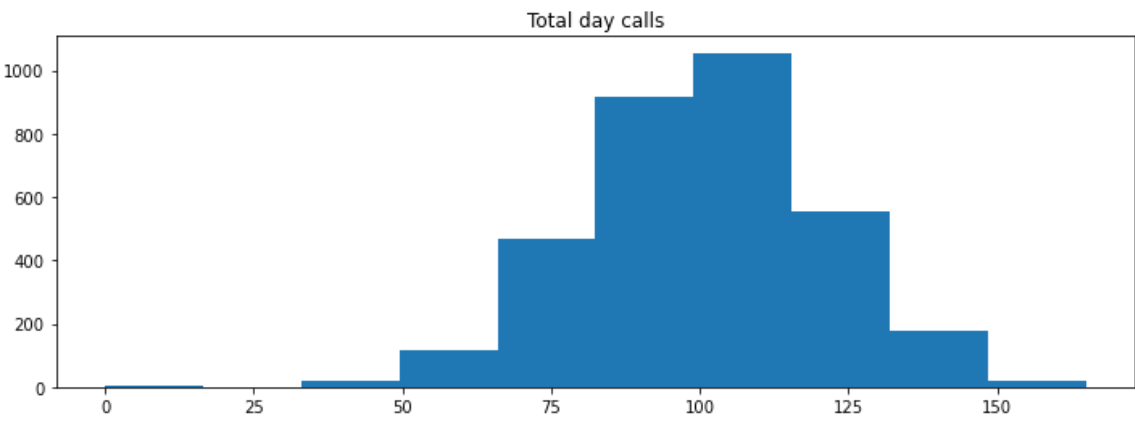



df.columns

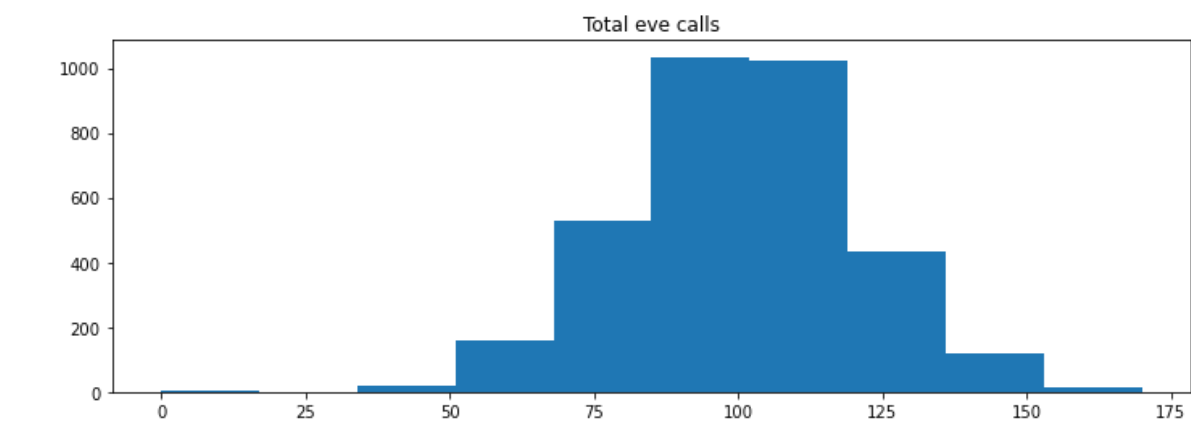

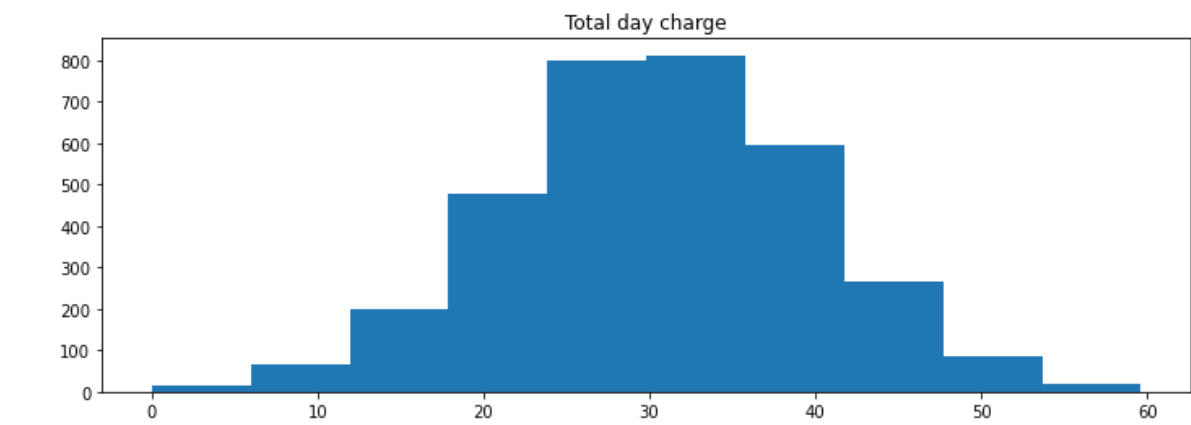
Out[24]:




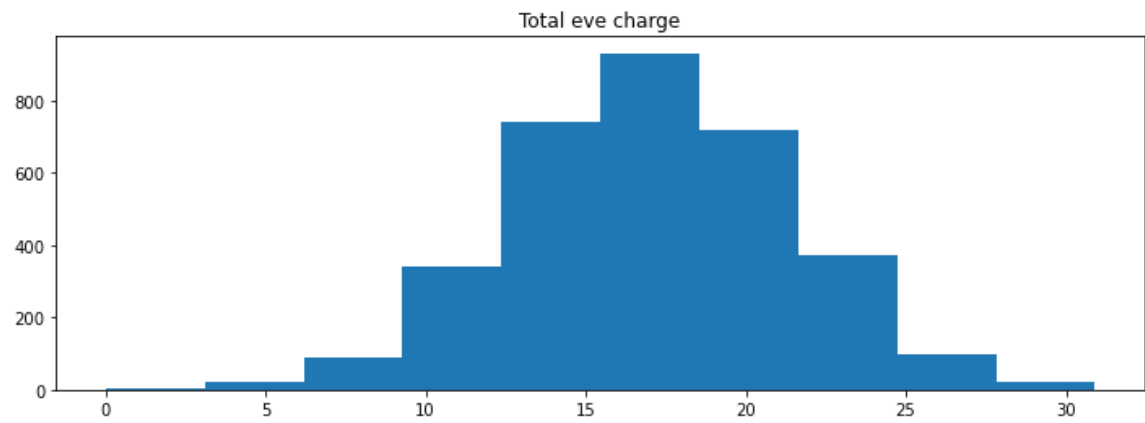
```
Index(['State', 'Account length', 'Area code', 'International plan',
      'Voice mail plan', 'Number vmail messages', 'Total day minutes',
      'Total day calls', 'Total day charge', 'Total eve minutes',
      'Total eve calls', 'Total eve charge', 'Total night minutes',
      'Total night calls', 'Total night charge', 'Total intl minutes',
      'Total intl calls', 'Total intl charge', 'Customer service calls',
      'Churn'],
      dtype='object')
```


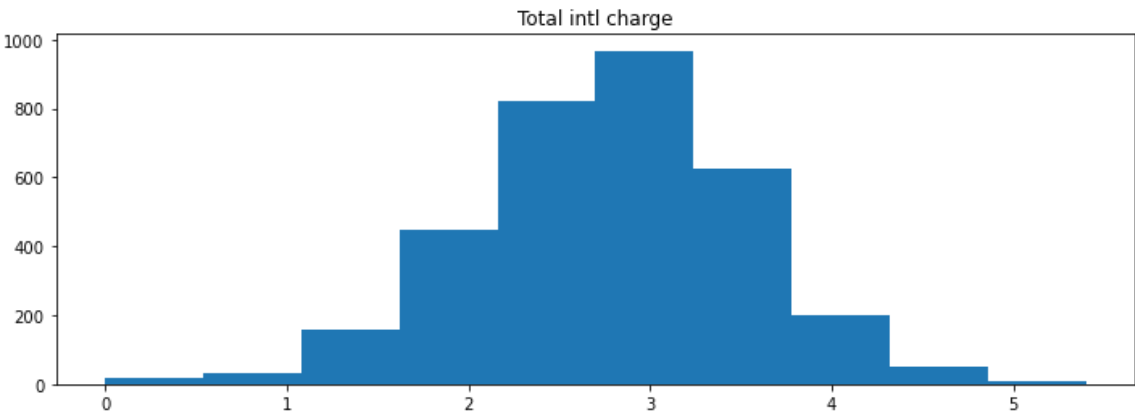

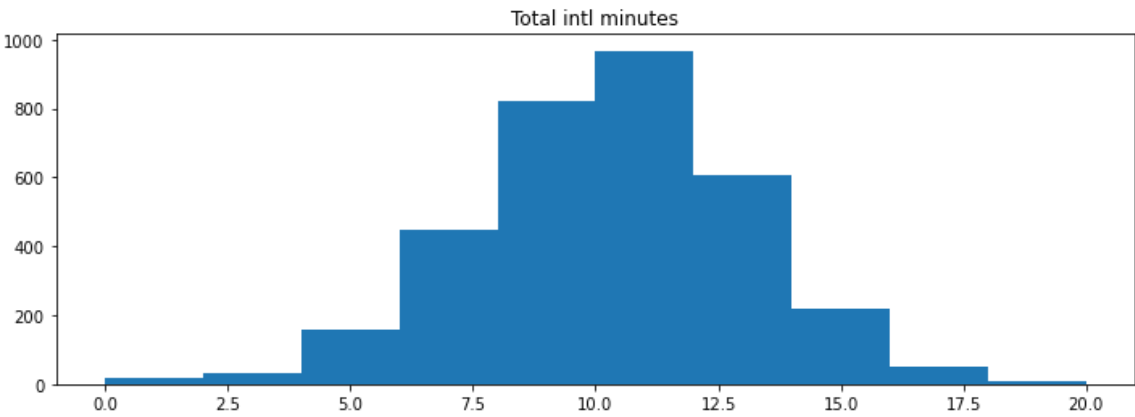
In [28]:

```
nums = ['Account length', 'Number vmail messages', 'Total day minutes',  
        'Total day calls', 'Total day charge', 'Total eve minutes',  
        'Total eve calls', 'Total eve charge', 'Total night minutes',  
        'Total night calls', 'Total night charge', 'Total intl minutes',  
        'Total intl calls', 'Total intl charge', 'Customer service calls']  
  
for col in nums:  
    plt.figure(figsize=(12,4))  
    plt.hist(df[col])  
    plt.title(col)  
    plt.show()
```









Observation -

- total xx minutes is having similar distribution to total xx charges
- Customer service calls, are having outliers present
- Number vmail messages is having multimodel distribution

Exploratory Analysis


- Correlation Analysis
- ANOVA
- Chi Square test

Correlation Analysis

In [32]:

```
corr = df.corr()

plt.figure(figsize=(16,10))
sns.heatmap(corr,annot=True,cmap='coolwarm')
plt.show()
```



Observation -

- slightly good correlating factors to customer churn include - Total day minutes and customer service calls
- Total xx minutes are exactly multiplier of total xx charge

ANOVA - Analysis of Variance

In [33]:

```
print(nums)
```

```
['Account length', 'Number vmail messages', 'Total day minutes', 'Total day calls', 'Total day charge', 'Total eve minutes', 'Total eve calls', 'Total eve charge', 'Total night minutes', 'Total night calls', 'Total night charge', 'Total intl minutes', 'Total intl calls', 'Total intl charge', 'Customer service calls']
```

In [35]:

```
xnum = df[nums]
y = df['Churn']

from sklearn.feature_selection import f_classif
fvalue, pvalue = f_classif(xnum,y)
```

In [36]:

```
for i in range(len(nums)):
    print(nums[i],pvalue[i])
```

```
Account length 0.33976000705720666
Number vmail messages 2.1175218402696038e-07
Total day minutes 5.300278227509361e-33
Total day calls 0.28670102402211844
Total day charge 5.30060595239102e-33
Total eve minutes 8.011338561256927e-08
Total eve calls 0.5941305829720491
Total eve charge 8.036524227754477e-08
Total night minutes 0.04046648463758881
Total night calls 0.7230277872081609
Total night charge 0.040451218769160205
Total intl minutes 8.05731126549437e-05
Total intl calls 0.002274701409850077
Total intl charge 8.018753583047257e-05
Customer service calls 3.900360240185746e-34
```

Observations -

- important informative features - Number vmail messages, total day mins, total eve mins, total night mins, total intl mins, total int calls, customer service calls

Chi Square test

In [39]:

```
cats = ['State', 'Area code', 'International plan', 'Voice mail plan']
```

In [41]:

```
xcat = df[cats]  
y = df['Churn']
```

In [42]:

```
from sklearn.preprocessing import LabelEncoder  
xcat['State'] = LabelEncoder().fit_transform(xcat['State'])  
xcat['International plan'] = LabelEncoder().fit_transform(xcat['International plan'])  
xcat['Voice mail plan'] = LabelEncoder().fit_transform(xcat['Voice mail plan'])
```

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:2: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

This is separate from the ipykernel package so we can avoid doing imports until

C:\ProgramData\Anaconda3\lib\site-packages\ipykernel_launcher.py:4: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
after removing the cwd from sys.path.

In [43]:

```
from sklearn.feature_selection import chi2  
chi_val, pvalue = chi2(xcat,y)
```

In [44]:

```
for i in range(len(cats)):
    print(cats[i],pvalue[i])
```

State 0.19214978695607624

Area code 0.4701527286099566

International plan 4.091734729415479e-46

Voice mail plan 5.28486023170551e-07

Observation -

- Important / informative features - International Plan, Voice mail plan




Data Visualization




In [46]:

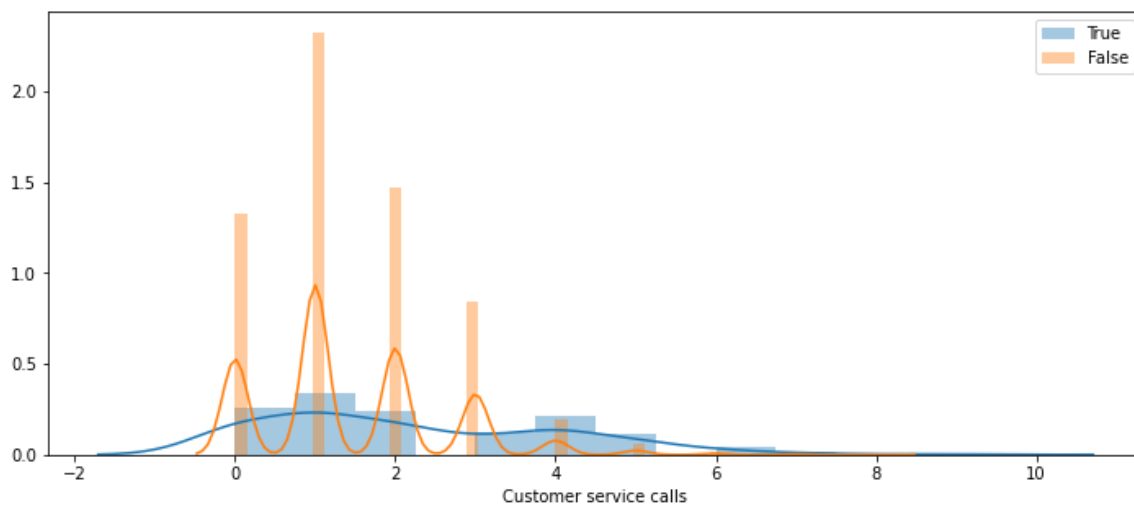
```
nums = ['Number vmail messages', 'Total day minutes', 'Total eve minutes', 'Total night m  
inutes',  
        'Total intl minutes', 'Total intl calls', 'Customer service calls']
```

In [47]:

```
for col in nums:
    plt.figure(figsize=(12,5))
    sns.distplot(df[col][df.Churn==True])
    sns.distplot(df[col][df.Churn==False])
    plt.legend([True,False])
    plt.show()
```







In [48]:

```
cats
```

Out[48]:

```
['State', 'Area code', 'International plan', 'Voice mail plan']
```


In [52]:

```
cats=['International plan', 'Voice mail plan']
for col in cats:
    pivot = pd.crosstab(df[col],df['Churn'],margins=True)
    print(pivot)
    ratio = pivot[True]/pivot['All']
    print(ratio)
    ratio.plot(kind='bar')
    plt.show()
```

Churn	False	True	All
International plan			
No	2664	346	3010
Yes	186	137	323
All	2850	483	3333


International plan

No 0.114950

Yes 0.424149

All 0.144914

dtype: float64



Churn	False	True	All
Voice mail plan			
No	2008	403	2411
Yes	842	80	922
All	2850	483	3333

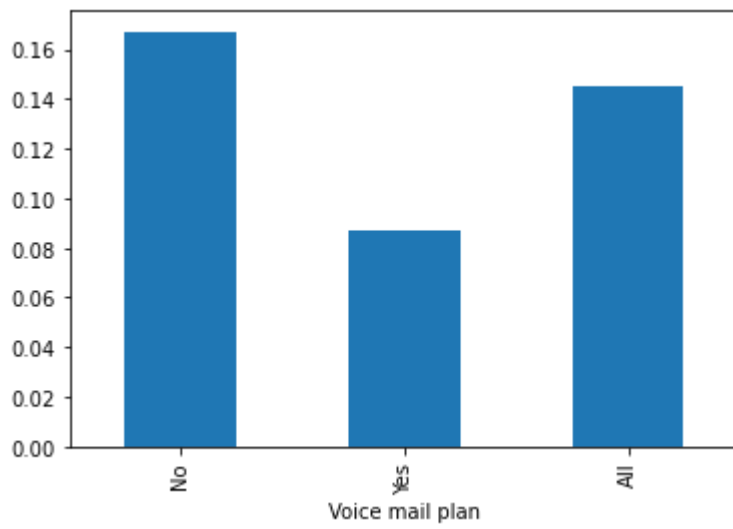
Voice mail plan

No 0.167151

Yes 0.086768

All 0.144914

dtype: float64



In []: